

System Development with Python: Week 8

Christopher Barker

UW Continuing Education

May 14, 2013

Table of Contents

- 1 Timing
- 2 Profiling
- 3 Performance Tuning

Performance Testing

“Premature optimization is the root of all evil”
– Donald Knuth

Profiling/timing

You can't optimize your code without knowing where the bottlenecks are.

Smarter people than me have said they they are almost always wrong when they try to logically determine where the slow code is. (I know I am)

... and how to speed it up

time.clock()

The really easy way:

```
import time

start = time.clock()
... do_some_stuff ...
print "It took %f seconds to run"%(time.clock - start)
```

It works, it's easy, and it gives a gross approximation

(use `time.clock()`, rather than `time.time()`)

timeit

The good way:

```
import timeit
```

```
timeit.timeit( statement, setup=some_stuff)
```

It's kind of a pain, but gives meaningful results.
(can also be called on the command line)

<http://docs.python.org/library/timeit.html>

(code/timing.py)

%timeit

The easy and good way:

ipython:

```
In [52]: import timing
```

```
In [53]: %timeit timing.primes_stupid(5)  
100000 loops, best of 3: 10.9 us per loop
```

Takes care of the setup/namespaces stuff for you

<http://ipython.org/ipython-doc/dev/interactive/tutorial.html>

Profiling

A profiler is a tool that describes the run time performance of a program, providing a variety of statistics

Helpful when you don't yet know where your bottlenecks are

The python profiler

```
python -m cProfile profile_example.py
```

spews some stats

<http://docs.python.org/library/profile.html>

python profiler

What you get:

ncalls the number of calls.

tottime the total time spent in the given function (and excluding time made in calls to sub-functions),

percall the quotient of tottime divided by ncalls

cumtime the total time spent in this and all subfunctions (from invocation till exit). This figure is accurate even for recursive functions.

percall the quotient of cumtime divided by primitive calls

(demo: `python -m cProfile profile_example.py`)

python profiler

You can also dump to a file:

```
$ python -m cProfile -o profile_dump profile_example.py
```

This gives you a binary file you can examine with pstats:

```
demo: $ python -m pstats
```

pstats

Running pstats

\$

\$ python -m pstats

Welcome to the profile statistics browser.

% read profile_dump

profile_dump% stats

Wed Aug 29 16:21:39 2012 profile_dump

51403 function calls in 0.032 seconds

Random listing order was used

ncalls	totttime	percall	cumtime	percall	filename:lineno(fu
51200	0.006	0.000	0.006	0.000	{method 'append' o
1	0.000	0.000	0.032	0.032	profile_example.py
1	0.001	0.001	0.032	0.032	profile_example.py
100	0.022	0.000	0.027	0.000	profile_example.py
100	0.002	0.000	0.024	0.000	profile_example.py

pstats commands

Commands:

- `help` help on pstats or particular command
- `stats` print the profile statistics
- `sort` sort by various data fields
- `strip` strips the leading path info from file names
- `callers` Print callers statistics
- `callees` Print callees statistics
- `quit` quits

Each has options to customize output

automating profile stats

cProfile and pstats are also modules

So you can script collection of profiles and stats

<http://docs.python.org/library/profile.html>

“Run Snake Run”

For a visual look at your profiling results:

`http://www.vrplumber.com/programming/runsnakerun/`

(pretty cool stuff!)

Performance Tips

Some common python performance issues:

`http:
//wiki.python.org/moin/PythonSpeed/PerformanceTips/`

(some nifty profiling tools described there, too)

LAB

Profiling lab

- run `timeit` on some code of yours (or `timing.py`, or..)
- run `iPython's %timeit` on the same code.
- try to make the factorial code in `timing.py` faster, and time the difference.
- write some code that tests one of the performance issues in:
`http://wiki.python.org/moin/PythonSpeed/PerformanceTips`
use one of the `timeits` to see if you can make a difference.
- try the profile tutorial at:
`http://pysnippet.blogspot.com/2009/12/profiling-your-python-code.html`

Wrap up

I hope you have an idea how to profile and time your code.

Try it on a part of your project

Next Week:

Student Presentations

Homework

Profile your project

Performance tune part of it

Get ready to present

Project Time!

Map out what you're going to present